

```

#1
import pandas as pd
import numpy as np
data = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
missing_values = data.isnull().sum()
print(missing_values)

satisfaction_level      0
last_evaluation          0
number_project           0
average_montly_hours    0
time_spend_company       0
Work_accident            0
left                      0
promotion_last_5years   0
sales                     0
salary                    0
dtype: int64

#2
data.head()
data
df = pd.DataFrame(data)
df.head()

      satisfaction_level  last_evaluation  number_project
average_montly_hours \
0                  0.38          0.53              2
157
1                  0.80          0.86              5
262
2                  0.11          0.88              7
272
3                  0.72          0.87              5
223
4                  0.37          0.52              2
159

      time_spend_company  Work_accident  left  promotion_last_5years
sales \
0                  3             0         1                  0
NaN
1                  6             0         1                  0
NaN
2                  4             0         1                  0
NaN
3                  5             0         1                  0
NaN
4                  3             0         1                  0
NaN

```

```
    salary
0      low
1  medium
2  medium
3      low
4      low

df.isna().sum()

satisfaction_level      0
last_evaluation         0
number_project          0
average_montly_hours   0
time_spend_company      0
Work_accident           0
left                     0
promotion_last_5years  0
sales                   0
salary                  0
dtype: int64

df["left"].unique()
array([], dtype=int64)

df["number_project"].unique()
array([], dtype=int64)

df.satisfaction_level.unique()
array([], dtype=float64)

df.last_evaluation.unique()
array([], dtype=float64)

df.time_spend_company.unique()
array([], dtype=int64)

df.Work_accident.unique()
array([], dtype=int64)

df.sales.unique()
array([], dtype=float64)
```

```

df.salary.unique()

array([], dtype=float64)

print(df.nunique())


Series([], dtype: float64)

print(df.dtypes)

Series([], dtype: object)

# Drop columns with only one unique value
df = df.loc[:, df.nunique() > 1]
df.corr()

Empty DataFrame
Columns: []
Index: []

import pandas as pd
import numpy as np
df = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
numeric_df = df.select_dtypes(include=['number'])

if numeric_df.empty:
    print("No numeric data found for correlation.")
else:
    correlation_matrix = numeric_df.corr()

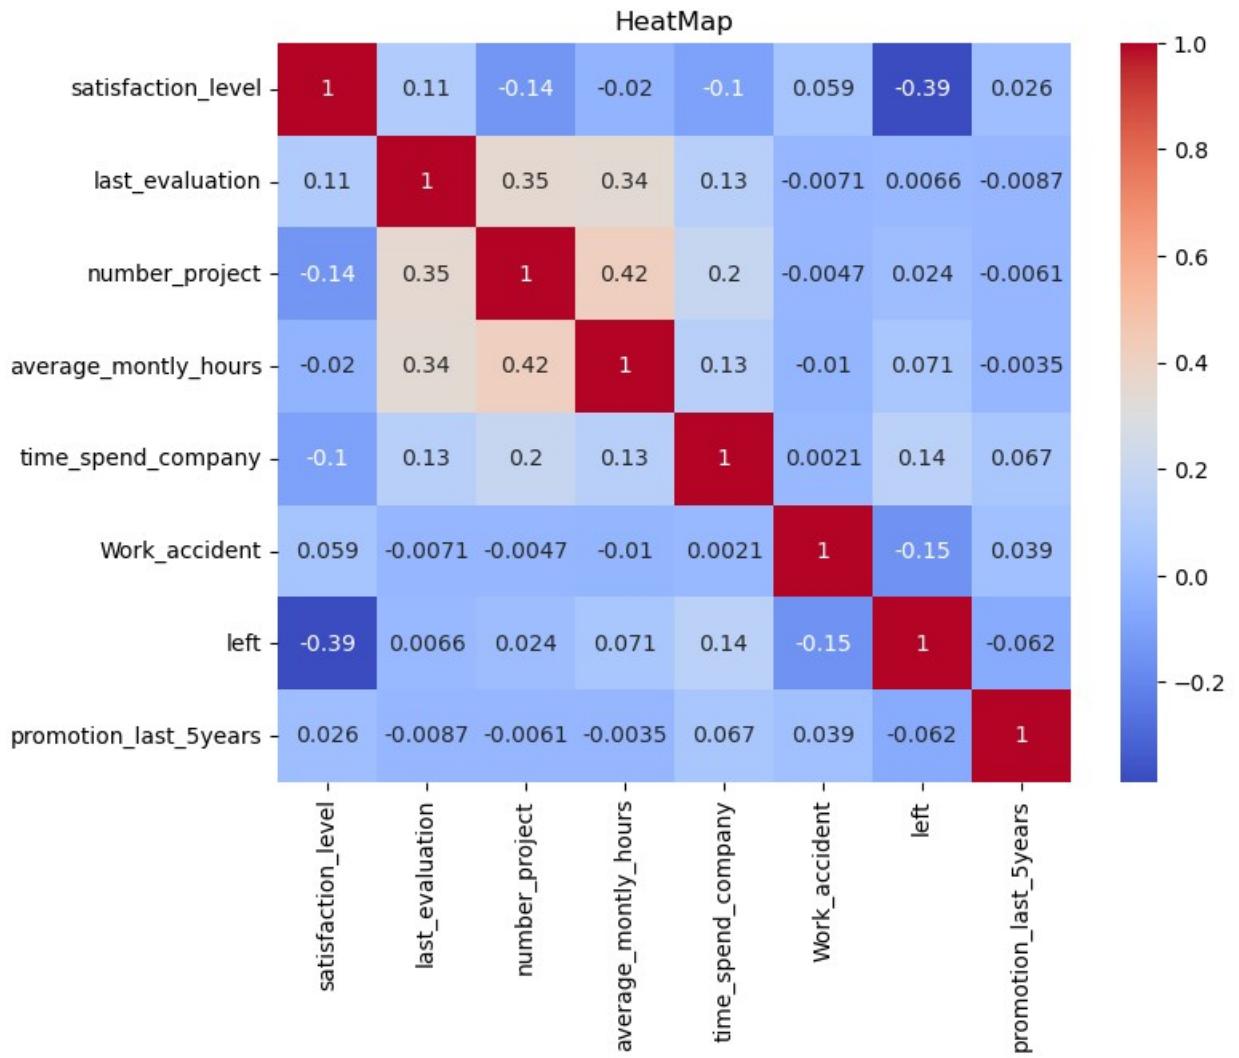
print(correlation_matrix)

            satisfaction_level  last_evaluation
number_project \
satisfaction_level           1.000000       0.105021      -
0.142970
last_evaluation                 0.105021       1.000000      -
0.349333
number_project                  -0.142970       0.349333      -
1.000000
average_montly_hours             -0.020048       0.339742      -
0.417211
time_spend_company                -0.100866       0.131591      -
0.196786
Work_accident                      0.058697      -0.007104      -
0.004741
left                            -0.388375       0.006567      -
0.023787
promotion_last_5years               0.025605      -0.008684      -
0.006064

```

	average_montly_hours	time_spend_company	\
satisfaction_level	-0.020048	-0.100866	
last_evaluation	0.339742	0.131591	
number_project	0.417211	0.196786	
average_montly_hours	1.000000	0.127755	
time_spend_company	0.127755	1.000000	
Work_accident	-0.010143	0.002120	
left	0.071287	0.144822	
promotion_last_5years	-0.003544	0.067433	
Work_accident left promotion_last_5years			
satisfaction_level	0.058697	-0.388375	0.025605
last_evaluation	-0.007104	0.006567	-0.008684
number_project	-0.004741	0.023787	-0.006064
average_montly_hours	-0.010143	0.071287	-0.003544
time_spend_company	0.002120	0.144822	0.067433
Work_accident	1.000000	-0.154622	0.039245
left	-0.154622	1.000000	-0.061788
promotion_last_5years	0.039245	-0.061788	1.000000

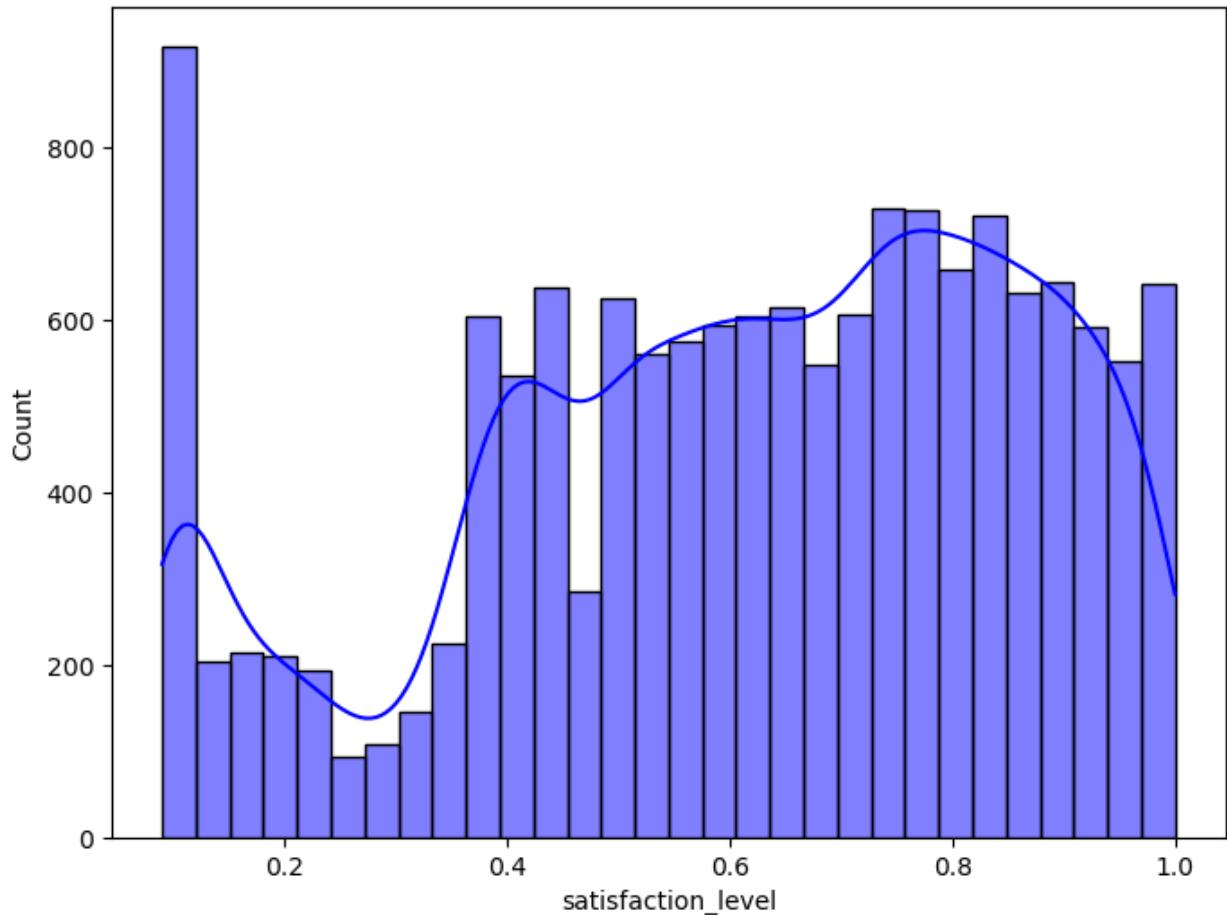
```
#heatmap
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize = (8,6))
sns.heatmap(correlation_matrix, annot = True,cmap ='coolwarm')
plt.title("HeatMap")
plt.show()
```



## #2.2 #Employee Satisfaction Distribution Plot

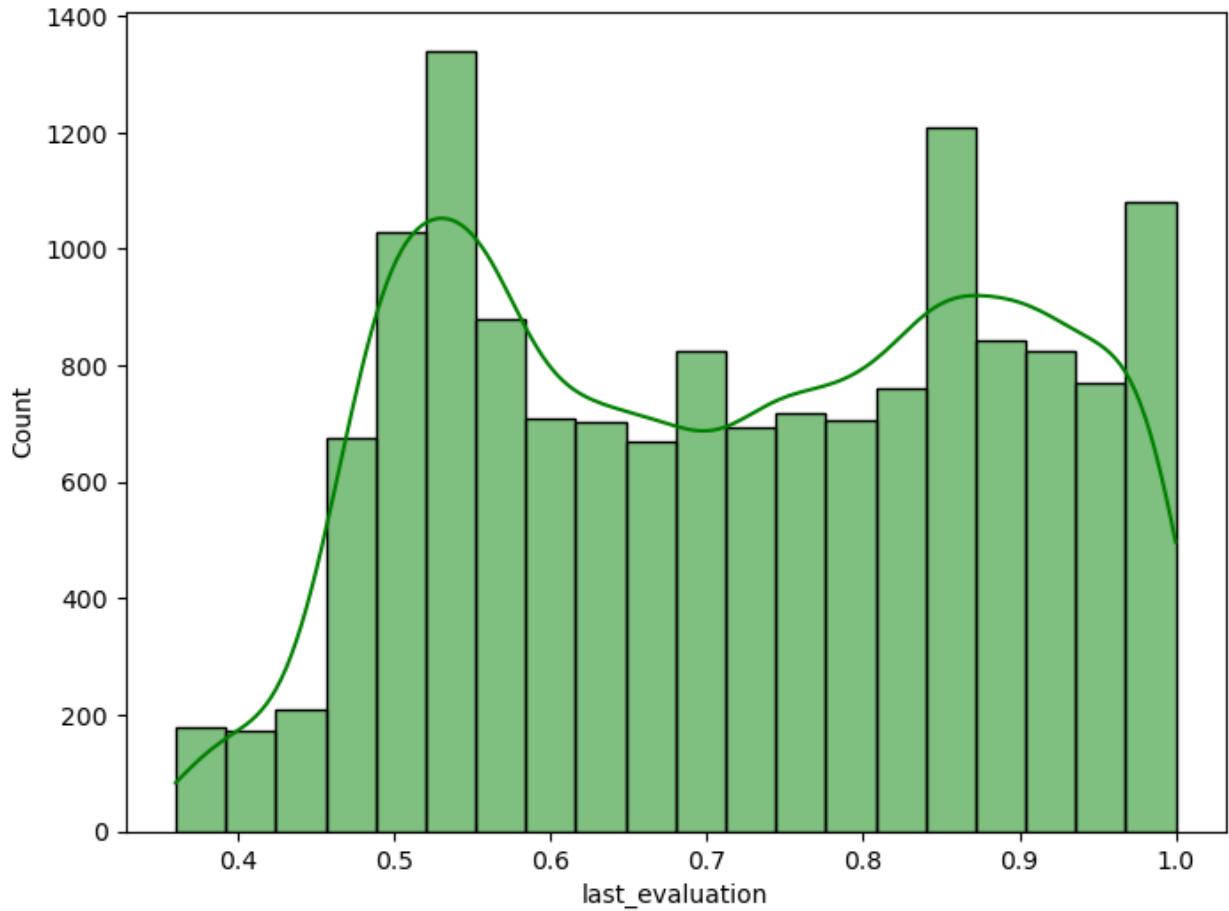
```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
plt.figure(figsize =(8,6))
satisfaction_data = df['satisfaction_level']
sns.histplot(satisfaction_data, kde = True , color = 'blue' , bins = 30)

<Axes: xlabel='satisfaction_level', ylabel='Count'>
```



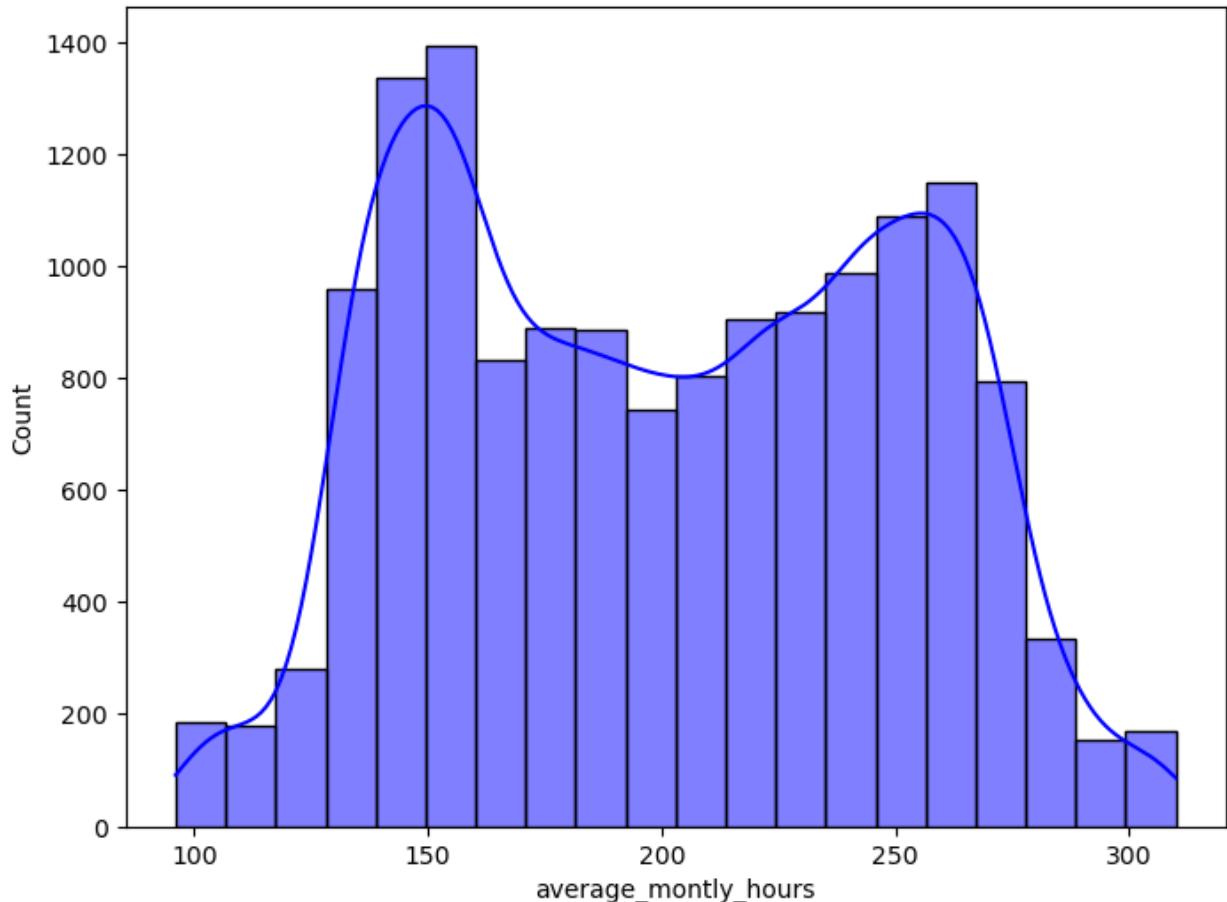
```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
plt.figure(figsize = (8,6))
last_evaluation = df['last_evaluation']
sns.histplot(last_evaluation, kde = True, color = 'green', bins = 20)

<Axes: xlabel='last_evaluation', ylabel='Count'>
```



```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
plt.figure(figsize = (8,6))
average_montly_hours = df['average_montly_hours']
sns.histplot(average_montly_hours, kde = True , color = 'blue', bins =20)

<Axes: xlabel='average_montly_hours', ylabel='Count'>
```



```

import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
data = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
left_employees = data[data['left'] == 1]
features = left_employees[['satisfaction_level', 'last_evaluation']]

# clustering
kmeans = KMeans(n_clusters=3, random_state=42)
clusters = kmeans.fit_predict(features)
left_employees['Cluster'] = clusters

# Visualize
plt.figure(figsize=(10, 6))
for cluster in range(3):
    cluster_data = left_employees[left_employees['Cluster'] == cluster]
    plt.scatter(
        cluster_data['satisfaction_level'],
        cluster_data['last_evaluation'],
        c='blue',
        s=100
    )

```

```

        label=f'Cluster {cluster}'
    )

plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1],
            color='black', marker='x', s=200, label='Centroids')
plt.title('Clustering of Employees Who Left')
plt.xlabel('Satisfaction Level')
plt.ylabel('Last Evaluation')
plt.legend()
plt.show()

# Print cluster centers
print("Cluster centers:")
print(kmeans.cluster_centers_)

```

```

/var/folders/r8/5p91n_mn2hj2fl6bdg5qg7100000gn/T/
ipykernel_86878/1034483085.py:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation:  
[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

left_employees['Cluster'] = clusters

```



```

Cluster centers:
[[0.41014545 0.51698182]
 [0.80851586 0.91170931]
 [0.11115466 0.86930085]]

#the cluster high satisfaction and low evaluation show inverse
relationship .

import pandas as pd
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
data = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
categorical_cols = data.select_dtypes(include=['object']).columns
numeric_cols = data.select_dtypes(include=['number']).columns

categorical_data = data[categorical_cols]
numeric_data = data[numeric_cols]

categorical_data_dummies = pd.get_dummies(categorical_data,
drop_first=True)
processed_data = pd.concat([numeric_data, categorical_data_dummies], axis=1)
X = processed_data.drop('left', axis=1) # Features (all columns
except 'left')
y = processed_data['left'] # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42, stratify=y)
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train,
y_train)
print("Class distribution before SMOTE:")
print(y_train.value_counts())
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_resampled).value_counts())

Class distribution before SMOTE:
left
0    7999
1    2500
Name: count, dtype: int64

Class distribution after SMOTE:
left
0    7999
1    7999
Name: count, dtype: int64

# After introduction smote the number , we see there was a
significant imbalance. It changed after the smote was applied .

```

```

#the number of minority class samples was increased.

import pandas as pd
from sklearn.model_selection import cross_val_predict
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.metrics import classification_report,
ConfusionMatrixDisplay
import matplotlib.pyplot as plt
data = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
X = data.drop(columns=["left"])
y = data["left"]
categorical_cols = ["sales", "salary"]
numeric_cols = [
    "satisfaction_level",
    "last_evaluation",
    "number_project",
    "average_montly_hours",
    "time_spend_company",
    "Work_accident",
    "promotion_last_5years",
]
# Preprocessing pipeline
preprocessor = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numeric_cols),
        ("cat", OneHotEncoder(), categorical_cols),
    ]
)
# Define models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42),
}
# Perform 5-fold cross-validation and generate reports
results = {}

for name, model in models.items():
    pipeline = Pipeline(steps=[("preprocessor", preprocessor),
    ("classifier", model)])

```

```

y_pred = cross_val_predict(pipeline, X, y, cv=5)
report = classification_report(y, y_pred, output_dict=True)
results[name] = report

# Display classification report and confusion matrix
print(f"\n{name} Classification Report:\n")
print(classification_report(y, y_pred))
ConfusionMatrixDisplay.from_predictions(y, y_pred)
plt.title(f"{name} Confusion Matrix")
plt.show()

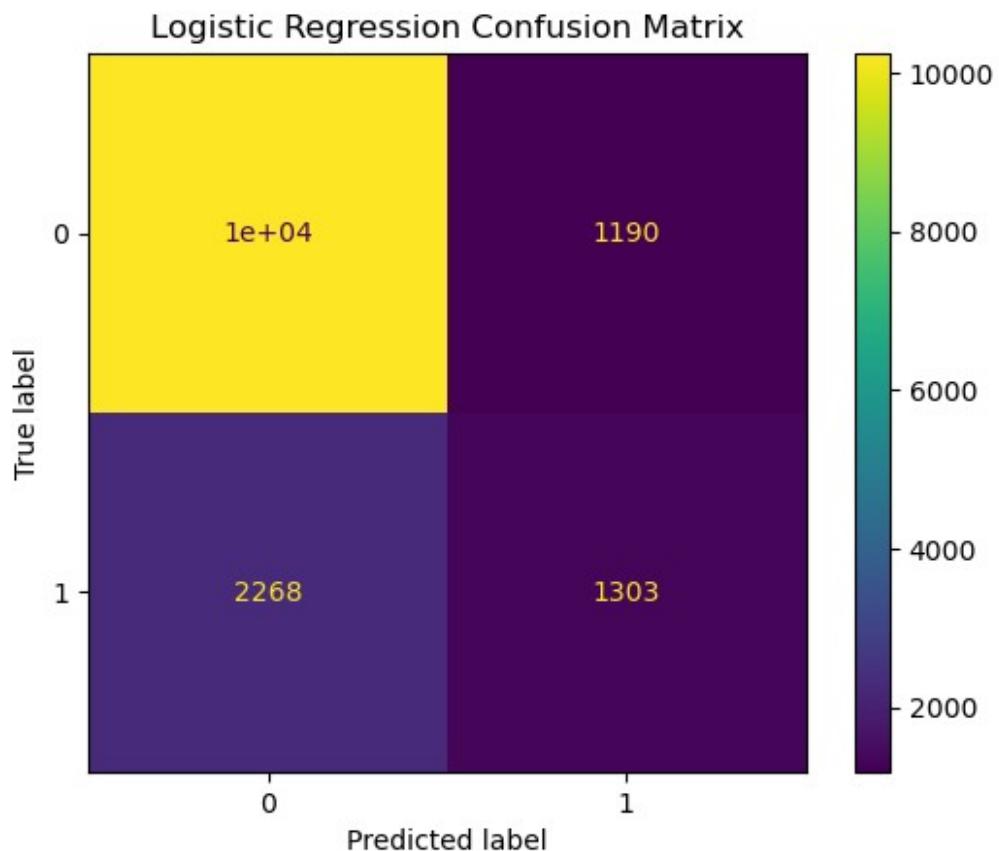
# Summary of results
results_summary = {
    name: {
        "Precision (1)": metrics["1"]["precision"],
        "Recall (1)": metrics["1"]["recall"],
        "F1-score (1)": metrics["1"]["f1-score"],
        "Accuracy": metrics["accuracy"],
    }
    for name, metrics in results.items()
}

# Display the summary
print("\nModel Performance Summary:\n")
summary_df = pd.DataFrame(results_summary).T
print(summary_df)

```

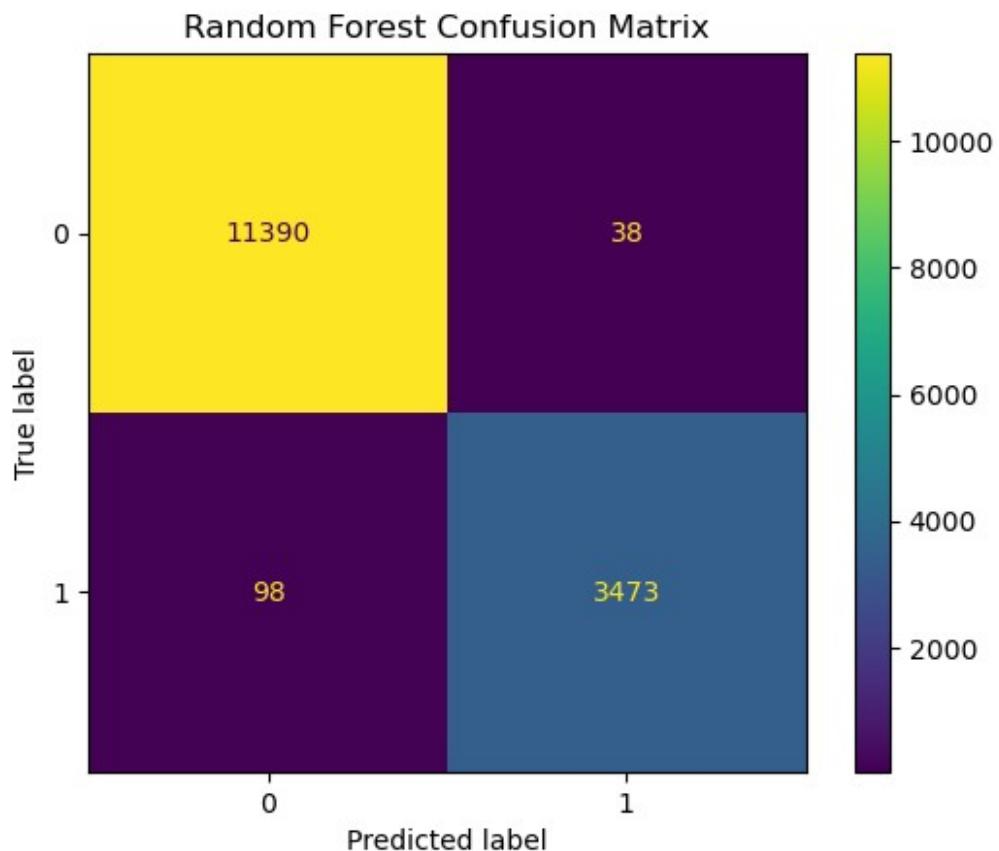
Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.82	0.90	0.86	11428
1	0.52	0.36	0.43	3571
accuracy			0.77	14999
macro avg	0.67	0.63	0.64	14999
weighted avg	0.75	0.77	0.75	14999



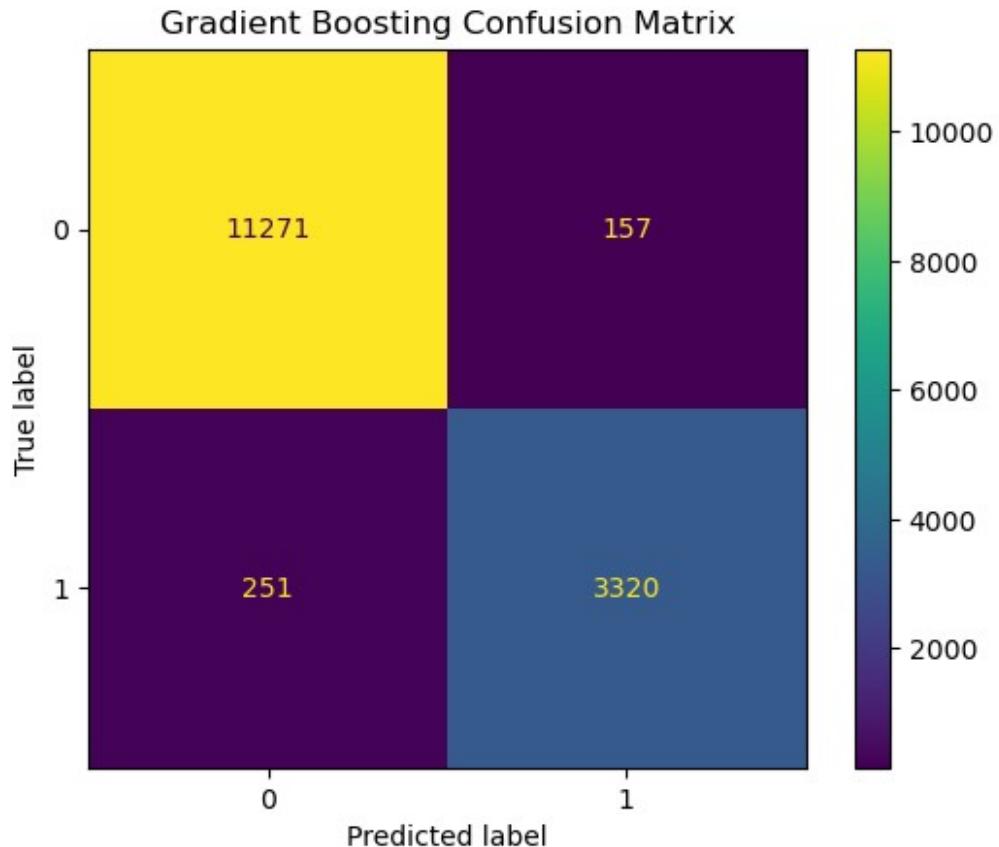
Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	11428
1	0.99	0.97	0.98	3571
accuracy			0.99	14999
macro avg	0.99	0.98	0.99	14999
weighted avg	0.99	0.99	0.99	14999



Gradient Boosting Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	11428
1	0.95	0.93	0.94	3571
accuracy			0.97	14999
macro avg	0.97	0.96	0.96	14999
weighted avg	0.97	0.97	0.97	14999



#### Model Performance Summary:

	Precision (1)	Recall (1)	F1-score (1)	Accuracy
Logistic Regression	0.522663	0.364884	0.429749	0.769451
Random Forest	0.989177	0.972557	0.980796	0.990933
Gradient Boosting	0.954846	0.929712	0.942111	0.972798

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.metrics import roc_curve, auc, confusion_matrix,
classification_report, RocCurveDisplay
import matplotlib.pyplot as plt
import seaborn as sns

data = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
X = data.drop(columns=["left"])

```

```

y = data["left"]

categorical_cols = ["sales", "salary"]
numeric_cols = [
    "satisfaction_level",
    "last_evaluation",
    "number_project",
    "average_montly_hours",
    "time_spend_company",
    "Work_accident",
    "promotion_last_5years",
]
preprocessor = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numeric_cols),
        ("cat", OneHotEncoder(), categorical_cols),
    ]
)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42),
}
# Train models, calculate ROC/AUC, and plot ROC curves
roc_auc_scores = {}
plt.figure(figsize=(10, 7))

for name, model in models.items():

    pipeline = Pipeline(steps=[("preprocessor", preprocessor),
    ("classifier", model)])

    pipeline.fit(X_train, y_train)

    y_pred_proba = pipeline.predict_proba(X_test)[:, 1]

    fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
    auc_score = auc(fpr, tpr)
    roc_auc_scores[name] = auc_score

```

```

plt.plot(fpr, tpr, label=f" {name} (AUC = {auc_score:.2f})")

# Plot details
plt.plot([0, 1], [0, 1], "k--", label="Random Guessing")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves for Models")
plt.legend(loc="lower right")
plt.show()

print("ROC/AUC Scores:")
for model, score in roc_auc_scores.items():
    print(f"{model}: AUC = {score:.3f}")

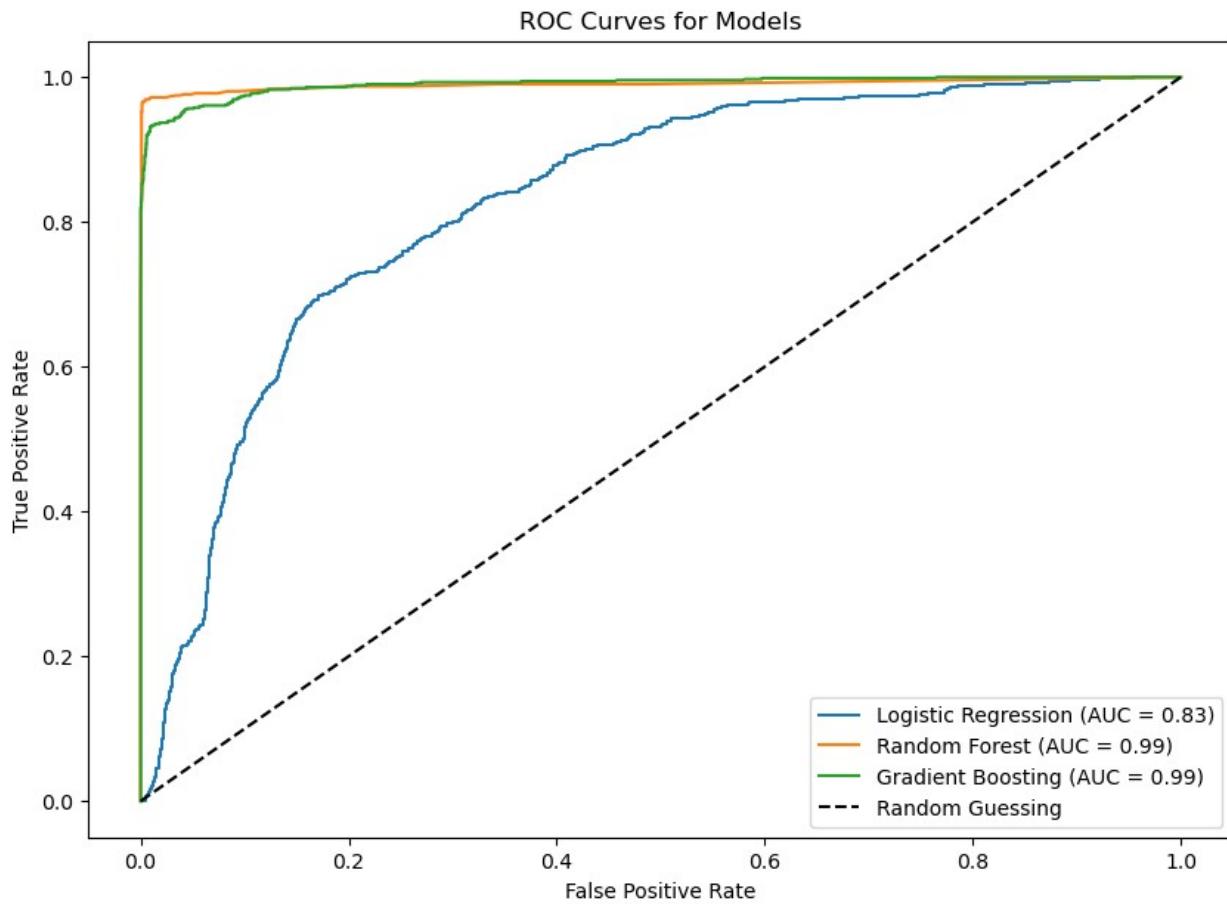
# Compute and Display confusion matrices and classification reports
for each model
for name, model in models.items():
    # Build a pipeline
    pipeline = Pipeline(steps=[("preprocessor", preprocessor),
    ("classifier", model)])

    # Train
    pipeline.fit(X_train, y_train)

    y_pred = pipeline.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)

    # Plot Confusion matrix
    plt.figure(figsize=(6, 5))
    plt.title(f"Confusion Matrix for {name}")
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
    xticklabels=["Stayed", "Left"], yticklabels=["Stayed", "Left"])
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
    print(f"Classification Report for {name}:")
    print(classification_report(y_test, y_pred))

```



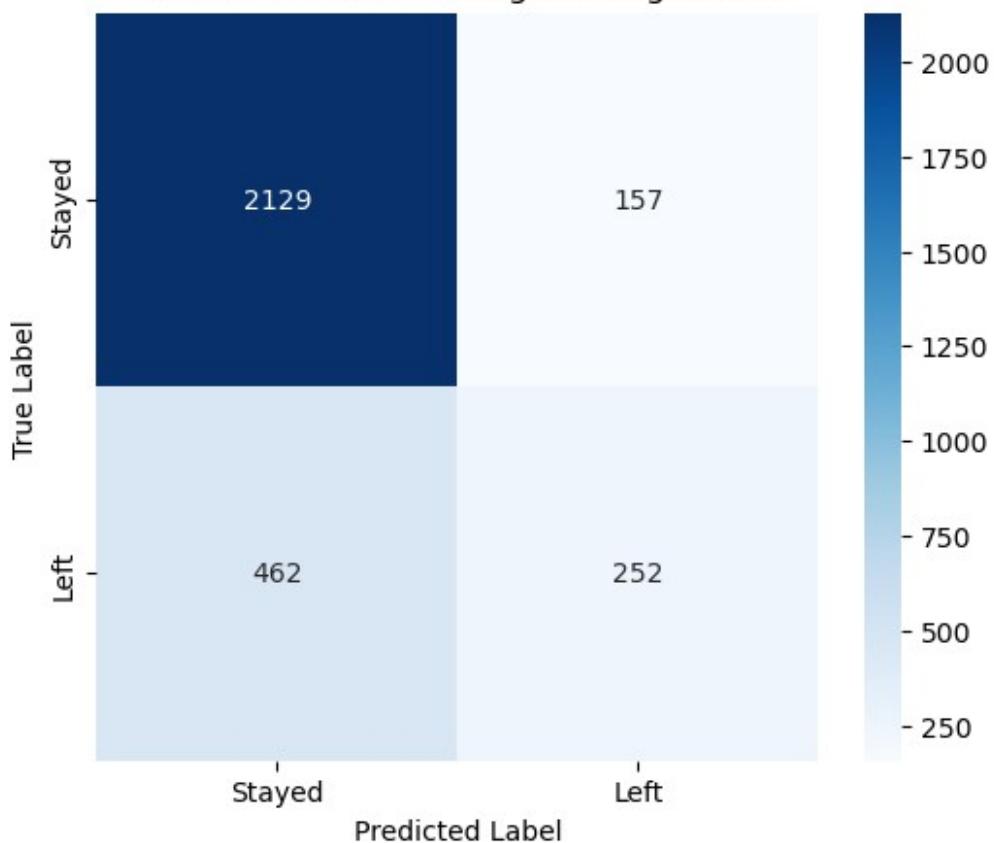
#### ROC/AUC Scores:

Logistic Regression: AUC = 0.828

Random Forest: AUC = 0.990

Gradient Boosting: AUC = 0.990

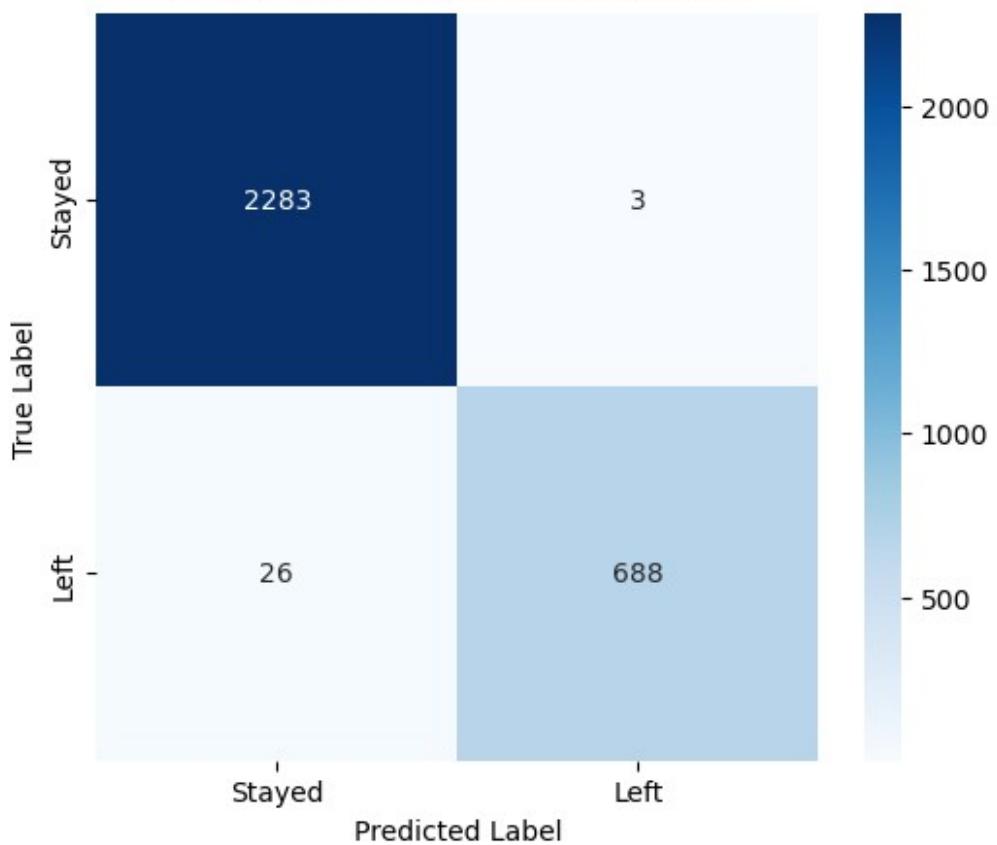
Confusion Matrix for Logistic Regression



Classification Report for Logistic Regression:

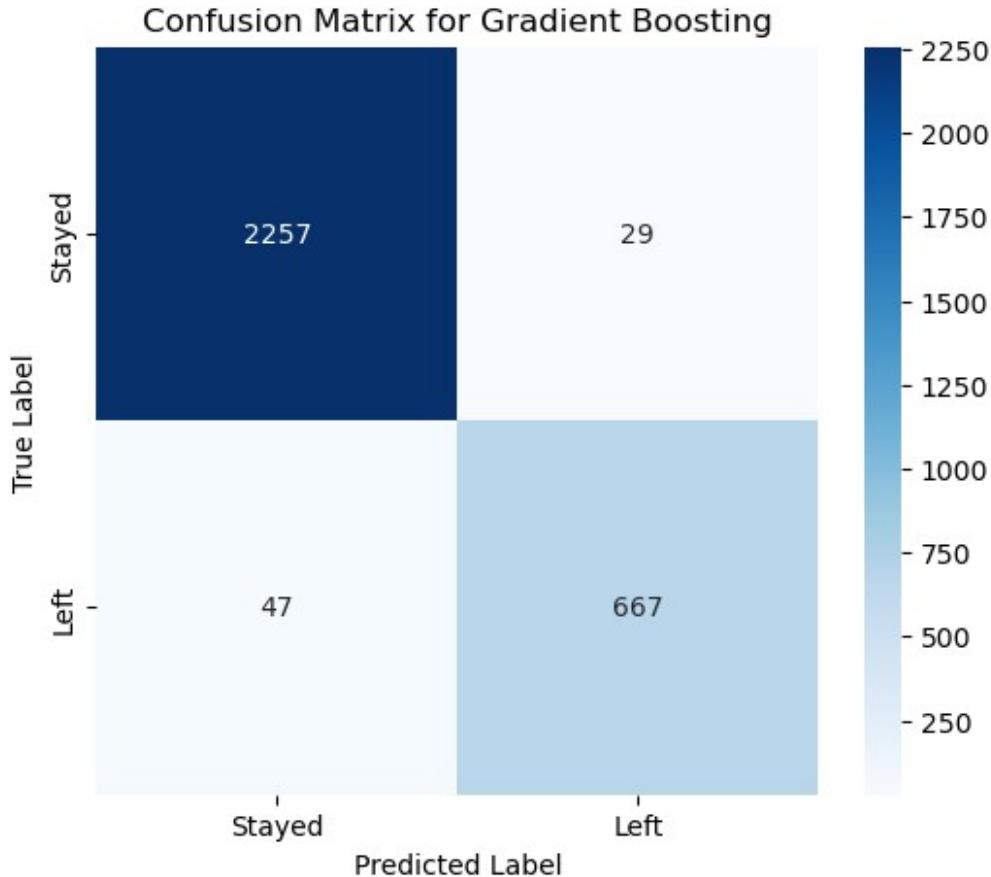
	precision	recall	f1-score	support
0	0.82	0.93	0.87	2286
1	0.62	0.35	0.45	714
accuracy			0.79	3000
macro avg	0.72	0.64	0.66	3000
weighted avg	0.77	0.79	0.77	3000

Confusion Matrix for Random Forest



Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	2286
1	1.00	0.96	0.98	714
accuracy			0.99	3000
macro avg	0.99	0.98	0.99	3000
weighted avg	0.99	0.99	0.99	3000



Classification Report for Gradient Boosting:				
	precision	recall	f1-score	support
0	0.98	0.99	0.98	2286
1	0.96	0.93	0.95	714
accuracy			0.97	3000
macro avg	0.97	0.96	0.96	3000
weighted avg	0.97	0.97	0.97	3000

#in my opinion , the random Forest achieves the best balance, with near-perfect precision, recall, and F1-score. It minimizes both false negative

#Using the best model, predict the probability of employee turnover in t

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

```

```

from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.metrics import roc_curve, auc, confusion_matrix,
classification_report, RocCurveDisplay
import matplotlib.pyplot as plt
import seaborn as sns
data = pd.read_csv(r"/Users/apple/Downloads/HR_comma_sep.csv")
X = data.drop(columns=["left"])
y = data["left"]
categorical_cols = ["sales", "salary"]
numeric_cols = [
    "satisfaction_level",
    "last_evaluation",
    "number_project",
    "average_montly_hours",
    "time_spend_company",
    "Work_accident",
    "promotion_last_5years",
]
preprocessor = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numeric_cols),
        ("cat", OneHotEncoder(), categorical_cols),
    ]
)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42),
}
best_model = RandomForestClassifier(random_state=42)
best_pipeline = Pipeline(steps=[("preprocessor", preprocessor),
("classifier", best_model)])
best_pipeline.fit(X_train, y_train)
y_pred_proba = best_pipeline.predict_proba(X_test)[:, 1]
zones = pd.DataFrame({
    "Employee": range(len(y_pred_proba)),
    "Probability": y_pred_proba
})
zones["Zone"] = pd.cut(
    zones["Probability"],
    bins=[0, 0.2, 0.6, 1],
    labels=["Safe Zone (Green)", "Low-Risk Zone (Yellow)", "High-Risk"
]

```

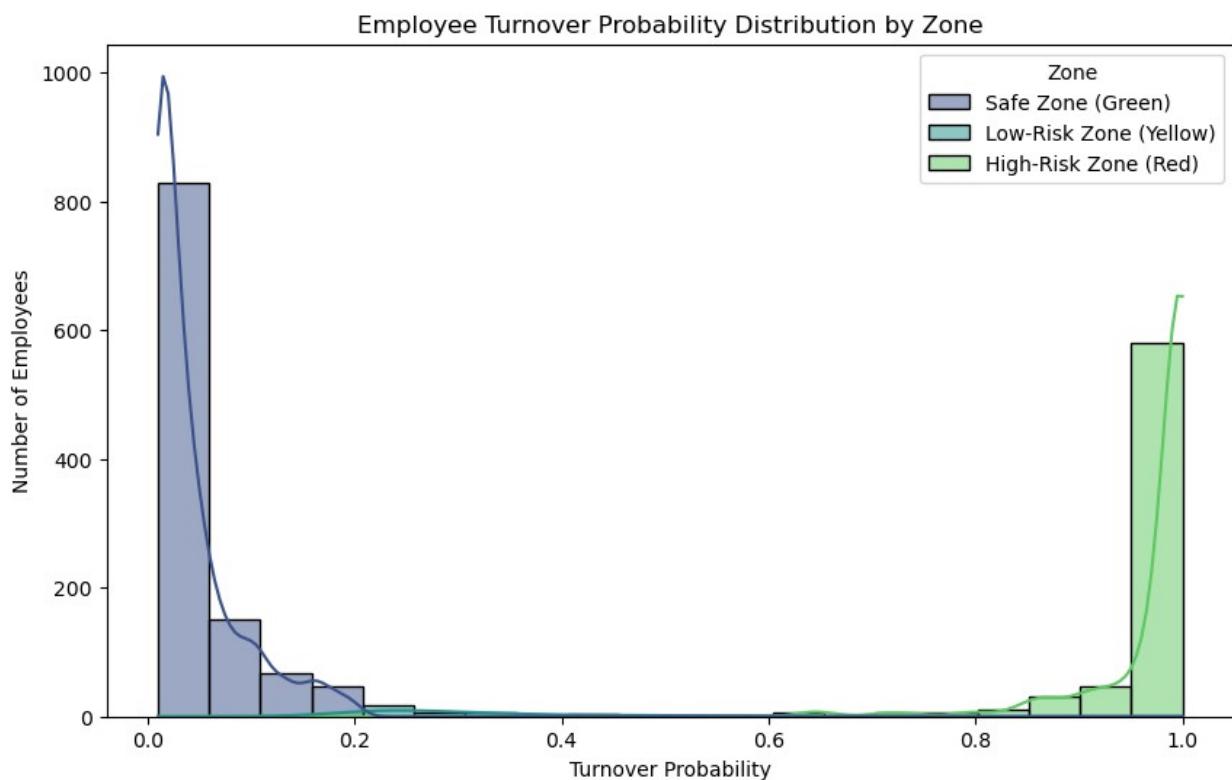
```

Zone (Red)"]
)
print("Categorized Zones:")
print(zones.head())
plt.figure(figsize=(10, 6))
sns.histplot(zones, x="Probability", hue="Zone", kde=True,
palette="viridis", bins=20)
plt.title("Employee Turnover Probability Distribution by Zone")
plt.xlabel("Turnover Probability")
plt.ylabel("Number of Employees")
plt.show()
print("\nRetention Strategies:")
print("1. Safe Zone (Green): Maintain current policies; these employees are highly satisfied.")
print("2. Low-Risk Zone (Yellow): Monitor satisfaction and engagement levels; small interventions may prevent turnover.")
print("3. High-Risk Zone (Red): Implement targeted retention strategies, such as personalized career growth plans or incentives.")

```

Categorized Zones:

Employee	Probability	Zone
0	0	0.04 Safe Zone (Green)
1	1	0.00 NaN
2	2	0.03 Safe Zone (Green)
3	3	0.00 NaN
4	4	0.00 NaN



**Retention Strategies:**

1. Safe Zone (Green): Maintain current policies; these employees are highly satisfied.
2. Low-Risk Zone (Yellow): Monitor satisfaction and engagement levels; small interventions may prevent turnover.
3. High-Risk Zone (Red): Implement targeted retention strategies, such as personalized career growth plans or incentives.